

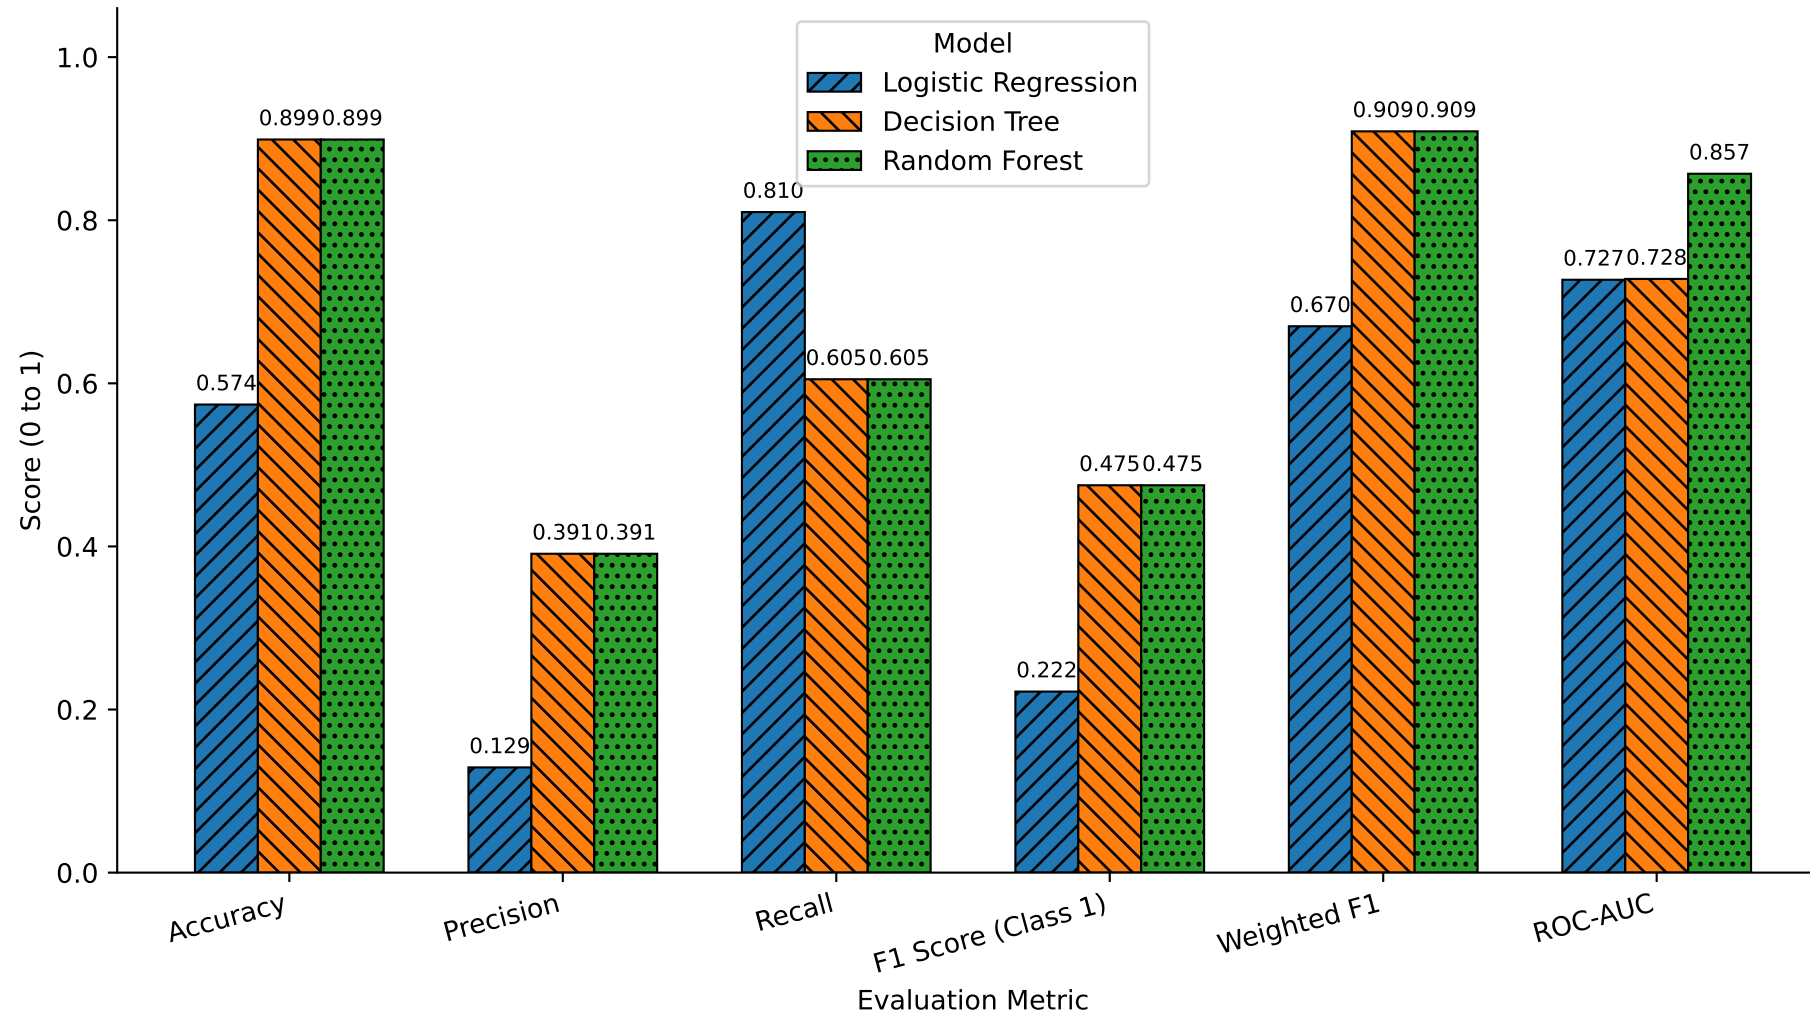


Figure 1. Predictive model performance on the held-out time-based test period.

Bars compare the three predictive models across Accuracy, Precision, Recall, F1 Score (Class 1), Weighted F1, and ROC-AUC. Patterned fills are used so the figure remains readable in grayscale or print.

Note. The Week-4 rule-defined reference is excluded from this chart because its perfect scores reflect label-definition agreement rather than independent predictive performance.